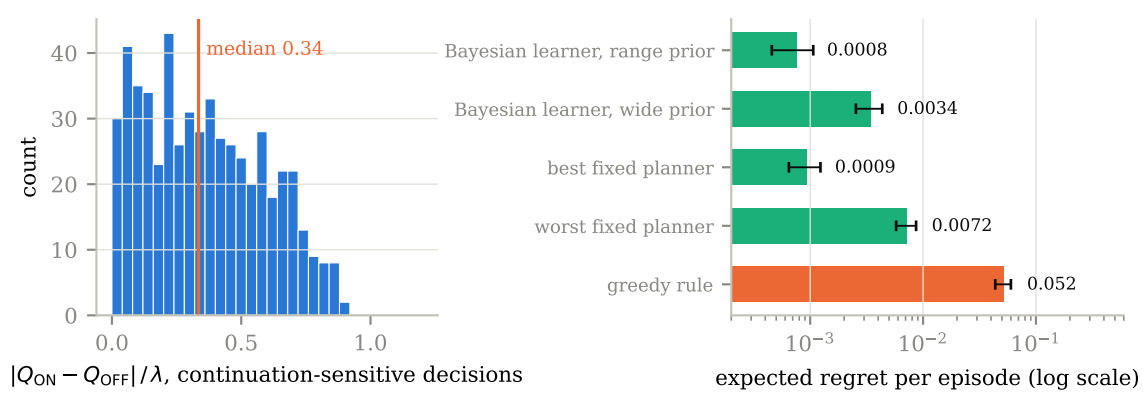
median 0.34
count
40
30
20
10
0
0.0
0.5
1.0
$|Q_{ON} - Q_{OFF}|/\lambda$, continuation-sensitive decisions
Bayesian learner, range prior
0.0008
Bayesian learner, wide prior
0.0034
best fixed planner
0.0009
worst fixed planner
0.0072
greedy rule
0.052
10^{-3}
10^{-2}
10^{-1}
expected regret per episode (log scale)